

# A Practical Theory for Symbolic Formula Discovery (Data— and Knowledge—Driven, LLM-Guided)

Cenk Alkan

October 6, 2025

## Abstract

I present a complete, implementation-ready theory for discovering concise and interpretable scientific formulas from data while enforcing domain knowledge. My design is LLM-guided but decisioned by verifiable criteria. I define: a protected numeric expression language and constant policy; a canonical complexity functional with a unique normal form; an interpretability score with a calibrated logistic map; a signed, normalized multi-objective loss; deterministic decoding with seed derivation; a sound unit/dimension type system; a deterministic mutual-information (MI) feature pipeline; a duplicate-aware population schedule with a correct 2D Pareto frontier; constant fitting procedures; a small MCTS for explanations; a synthetic-data protocol preserving identifiability; and an acceptance harness that certifies equality by a symbolic zero-FN route, with refutation on a rational lattice and a protected floating probe. Every piece is specified so that implementation reduces to faithfully encoding these rules.

## 1 Problem Statement and Objectives

Given inputs  $X \in \mathbb{R}^{n \times d}$  with variable meanings and SI dimensions, and a target  $y \in \mathbb{R}^n$  (or labels), I search for  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  that balances accuracy, simplicity, and interpretability under unit consistency. LLMs generate candidates; I select using formal criteria:

$$\min_f (E(f), C(f), -S(f)),$$

where  $E$  is a normalized error,  $C$  is my complexity, and  $S \in (0, 1)$  is a calibrated interpretability score. I report a 2D Pareto frontier in  $(C, E)$  and use a scalarized loss involving  $S$  only for survivor ranking.

## 2 Expression Language and Protected Numerics

### 2.1 Primitive set and guards

I restrict evaluation to a total, elementwise set of primitives with explicit guards:

$$\begin{aligned} \text{padd}(x, y) &= x + y, & \text{psub}(x, y) &= x - y, & \text{pmul}(x, y) &= xy, \\ \text{pdiv}(x, y) &= \frac{x}{y^\star}, & y^\star &= \text{sign}(y) \max(|y|, \varepsilon), \\ \text{pneg}(x) &= -x, & \text{pabs}(x) &= |x|, & \text{plog}(x) &= \log(|x| + \varepsilon_{\log}), \\ \text{pexp}(x) &= \exp(\text{clip}(x, -C_{\text{exp}}, C_{\text{exp}})), & \text{psqrt}(x) &= \sqrt{\max(x, 0)}, \\ \text{psin}, \text{pcos}, \text{ptanh}, & \text{ and fixed powers } p \in \{-3, -2, -1, -\tfrac{1}{2}, \tfrac{1}{2}, 2, 3\}. \end{aligned}$$

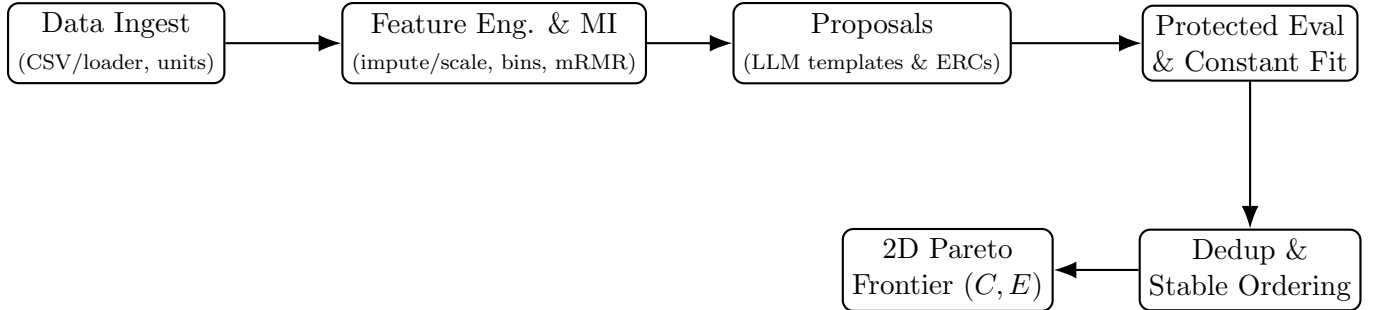


Figure 1: Canonical end-to-end process: data ingest  $\rightarrow$  feature/MI  $\rightarrow$  proposals  $\rightarrow$  protected evaluation  $\Downarrow$  dedup/ordering  $\Leftarrow$  2D Pareto frontier  $\rightarrow$  acceptance (symbolic  $\rightarrow$  lattice  $\rightarrow$  `float_probe`)  $\rightarrow$  logging/artifacts.

I treat  $\{\pi, e\}$  as read-only constants. Ephemeral random constants (ERCs) appear as terminals and may be refit after structure discovery.

**Proposition 1** (Well-defined objective). *Any expression  $f$  built from my protected primitives is finite and Borel-measurable on  $\mathbb{R}^m$ . If inputs are absolutely continuous and  $y$  is finite a.s., sample losses (MSE/MAE) are finite a.s. and well-defined.*

## 2.2 Grammar and sandbox

I use a minimal, effect-free grammar:  $\{+, -, \times, \text{pdiv}, \text{plog}, \text{pexp}, \text{psqrt}, \text{psin}, \text{pcos}, \text{ptanh}\}$  plus the fixed powers and constants/ERCs. Static checks enforce: (i) one function body, no imports or I/O; (ii) no builtins; (iii) calls only to protected numerics. Runtime executes in a fresh process with empty builtins and a restricted numeric proxy.

**Proposition 2** (Safety). *Under this whitelist and sandbox, any accepted artifact can only evaluate the intended numeric expression; file/network/process effects and capability escalation are ruled out by syntax and the execution environment.*

## 3 Constant Policy and Fitting

I sample ERCs from a mixture: discrete mass on small rationals and  $\{\pm\pi, \pm e\}$ , a uniform component on  $[-3, 3]$ , and a signed log-uniform component. After structure discovery, I fit constants by least squares: closed-form for linear-in-parameter templates; otherwise trust-region reflective least squares with deterministic multi-starts.

**Proposition 3** (Linear recovery). *For noiseless  $y = \Phi\theta^*$  with full column rank, least squares returns  $\hat{\theta} = \theta^*$ . With  $y = \Phi\theta^* + \varepsilon$  and  $\mathbb{E}\varepsilon = 0$ ,  $\hat{\theta} \rightarrow \theta^*$  as  $\text{Var}(\varepsilon) \rightarrow 0$ .*

## 4 Complexity Functional $C$

### 4.1 Definition and canonicalization

I parse to a canonical form and score nodes as follows. Leaves: variables cost 1; constants cost 0 for  $\{0, 1\}$ , cost 1 for  $\{2, -1, \frac{1}{2}, \pi, e\}$ , else digit-length proxy (floats simplified when exact). Internal nodes: Add/Mul cost  $(k - 1) + \sum C(\cdot)$ ; small rational/integer Pow costs  $2 + C(\text{base})$ ; other numeric exponent costs  $3 + C(\text{base})$ ; symbolic exponent costs  $4 + C(\text{base}) + C(\text{exp})$ ;  $\{\log, \exp, \sin, \cos, \tanh\}$  cost  $3 + \sum C$ ; Abs costs  $2 + \sum C$ . Canonicalization is an explicit, finite, terminating transform set: flatten/sort Add/Mul, remove neutrals, normalize numeric factors.

**Lemma 1** (Unique normal form). *My transform set is semantics-preserving, terminates, and yields a unique normal form for any input expression.*

**Definition 1** (Normalized complexity).  $C_{\min}(e) = C(\text{canon}(e))$  where canon applies my fixed transform sequence.

**Proposition 4** (Subadditivity). *For  $m$ -ary  $f$  and expressions  $g_1, \dots, g_m$ , I have  $C(f(g_1, \dots, g_m)) \leq C(f) + \sum_j C(g_j)$ .*

## 5 Interpretability Score $S$ and Calibration

I grade five criteria in  $[-20, 20]$ : (1) structural simplicity, (2) variable semantics/units, (3) modularity, (4) parameter interpretability, (5) domain behavior. The raw total  $x \in [-100, 100]$  is mapped to  $S(x) = \sigma(a(x - b))$  with  $a \leq 0.02$  so the global Lipschitz constant is  $L = a/4 \leq 0.005$ .

**Proposition 5** (Monotone and Lipschitz).  *$S$  is strictly increasing and globally Lipschitz with constant  $a/4$ .*

## 6 Loss, Normalization, and Ranking

### 6.1 Metrics and normalization

For regression,  $E = \text{clip}(1 - R^2, 0, 1)$  with a bounded NMAE fallback if  $R^2$  is undefined. For classification, I use cross-entropy divided by  $\log K$ . I normalize  $E, C, S$  by robust min-max within a fixed batch snapshot. Let  $N_E, N_C, N_S$  denote these snapshot normalizations of  $E, C, S$ , respectively.

**Definition 2** (Scalarized loss). *With  $(\alpha, \beta, \gamma) \geq 0$ ,*

$$L = \alpha N_E(E) + \beta N_C(C) - \gamma N_S(S).$$

**Proposition 6** (Monotone in  $S$  per snapshot). *For a fixed normalization snapshot,  $L$  is strictly nonincreasing in  $S$  when  $\gamma > 0$ .*

## 7 Units and Dimensional Typing

### 7.1 Model and rules

I model dimensions as exponent vectors  $d \in \mathbb{Z}^7$  over  $(M, L, T, I, \Theta, N, J)$ . Multiplication/division add/subtract exponents; rational powers scale exponents;  $\exp, \log, \sin, \cos, \tanh, \sqrt{\cdot}$  require dimensionless inputs; Abs preserves the input dimension.

**Proposition 7** (Soundness). *Given an environment  $\Gamma : x \mapsto d_x$ , my checker either returns a dimension  $d$  (derivation  $\Gamma \vdash e : d$ ) or rejects with a specific rule violation. Any accepted expression is dimensionally consistent.*

## 8 Feature Engineering and MI Selection

My pipeline: train-only mean-imputation; standard scaling; engineered transforms (polynomial, trigs,  $\log(|x| + \varepsilon)$ , pairwise products); plug-in MI  $I_b$  with equal-width bins ( $b = 16$ ); a stable top- $K$  prefilter and a greedy mRMR-style diversity selection with a fixed correlation cutoff.

**Proposition 8** (Consistency for discretized MI). *With fixed finite partitions and i.i.d. samples, the plug-in  $I_b(X; Y)$  converges almost surely to  $I(X_b; Y_b)$ , the MI between discretized variables.*

## 9 Population Schedule, Deduplication, and Budget

I use a steady-state loop with typical settings  $N=200$  initial, survivor cap  $K=30$ , new per round  $J=10$ , horizon  $T=500$ . I control duplicates by canonical structural keys, numeric fingerprints on a fixed grid, and (when needed) a symbolic certificate; survivors are ordered by  $(L, \text{canonical form})$  stably.

**Proposition 9** (Partial order). *Define  $f \preceq g$  if either their canonical forms match, or numeric fingerprints match and  $\text{simplify}(f - g) = 0$  and  $C_{\min}(f) \leq C_{\min}(g)$ . Then  $\preceq$  is a partial order; antisymmetry holds on canonical representatives.*

## 10 Pareto Frontier in Two Dimensions

For  $(C, E)$  minimization with tolerance  $\varepsilon \geq 0$ , I sort stably by  $(C, E)$  and sweep left-to-right, accepting a point iff  $E < E^* - \varepsilon$  (then updating  $E^*$ ). This returns exactly the  $\varepsilon$ -nondominated set in  $O(n \log n)$ .

**Proposition 10** (Correctness). *The sweep returns all and only  $\varepsilon$ -nondominated points (with stable tie handling for duplicates).*

## 11 Splits, Seeds, and Leakage Control

I use seeded stratified holdouts and  $K$ -folds. For grouped entities, I ensure group-disjoint splits. All partitions are deterministic in the seed, and my group-aware procedures eliminate entity leakage by construction.

**Proposition 11** (Stratified holdout bias). *Under exchangeability and classwise sampling without replacement, the stratified holdout estimate equals the class mixture of class risks; rounding effects induce  $O(1/N)$  bias.*

## 12 Deterministic Decoding and Prompt Discipline

### 12.1 Backend, endpoint, and model

Unless otherwise noted, all completions are generated by **llama3.1** served through the **Ollama** HTTP API at <http://127.0.0.1:11434>. The implementation treats this endpoint as the default

local backend; if an environment override is provided, it is recorded in the run manifest. When available, the model digest reported by Ollama is persisted alongside `model_id` to pin the exact binary.

## 12.2 Seed derivation and knobs

I fix decoding knobs (temperature, top- $k$ , top- $p$ , repetition penalty, token budget) and derive a 64-bit seed by

$$\text{seed}(s, p, i, a) = \text{uint64}(\text{SHA256}(s\|p\|i\|a)),$$

where  $s$  is the session ID,  $p$  a prompt tag,  $i$  an item index, and  $a$  a retry counter. Under *llama3.1* via Ollama, these knobs and the seed are passed deterministically through the client so that, given identical prompts and the same model digest, the sampled completions are reproducible.

**Proposition 12** (Determinism). *If base logits are deterministic in  $(p, y_{<t})$  at the backend and the only randomness comes from my seeded PRNG, completions are deterministic in  $(s, p, i)$  (and retry schedule).*

## 12.3 Prompt contract

I constrain LLM outputs to a single fenced Python function `def f(env):` with no imports and only whitelisted names. I enforce stop sequences so I always extract exactly one artifact per sample.

## 13 Explanation Search via MCTS

I include a small MCTS with UCB1 selection ( $c = \sqrt{2}$ ), bounded rollouts, and fixed budgets to generate human-facing explanations. Rewards are bounded in  $[0, 1]$  and distinct from  $S$ .

**Remark 1** (Use in practice). *I use UCB1 as a robust exploration–exploitation policy; in nonstationary trees it is a practical heuristic, not a formal optimality claim.*

## 14 Synthetic Data and Identifiability

For rediscovery tests, I sample over safe domains, affine-normalize features into  $[0, 1]^d$ , and (optionally) add Gaussian noise to normalized targets. Success means the planted form lies on my  $(C, E)$  frontier.

**Proposition 13** (Normalization preserves class). *If  $u_j = (x_j - L_j)/(U_j - L_j)$  are strictly increasing affine maps and the hypothesis class is closed under such reparameterizations, normalization preserves representability and structural identifiability (constants rescale).*

## 15 Acceptance Criteria and Certification

I canonicalize  $f$  (truth) and  $g$  (candidate) and apply a three-stage contract:

1. **Symbolic certificate (zero-FN).** If  $\text{simplify}(f - g) = 0$ , I accept with method `symbolic`.
2. **Deterministic rational-lattice refutation.** If symbolic fails, I search a fixed rational lattice in exact arithmetic. A finite, non-singular witness implies  $f \neq g$  and I reject with method `reject`.

3. **Protected floating probe.** If the lattice finds no witness, I probe on a safe float grid at tight tolerance and compute a simple miss-probability bound. Passing probes accept with method `float_probe`.

**Proposition 14** (Determinism and soundness). *For fixed  $(f, g)$ , variable order, and horizon, the lattice is deterministic; any returned witness is a certified counterexample. The symbolic route has zero false negatives. Absent both, a float probe with  $m$  i.i.d. points misses a region of disagreement of relative measure  $\delta > 0$  with probability at most  $e^{-\delta m}$ .*

## 16 Reproducibility and Logging

I record a manifest (decoding knobs, seed, environment) and canonical JSONL events. The manifest includes the LLM binding:

```
{model_id = "llama3.1", backend = "ollama",
  endpoint = "http://127.0.0.1:11434",
  model_digest = (if available)}.
```

The manifest hash is the SHA-256 of a canonical JSON core; event timestamps are logical and monotone. For each accepted/rejected pair I persist a compact artifact (`accept_proof.json`) and the canonical forms (`forms.txt`).

## 17 Scope and Limitations

Token and time budgets limit large-scale exploration. My acceptance is deliberately conservative: exact symbolic equalities are decisive; the lattice is refutation-only; the float probe is tightly guarded. I can layer broader proposal sets, hierarchical search, or retrieval without changing these foundations.

## Conclusion

I fixed a protected numeric language and constant policy, a canonical complexity, a calibrated interpretability score, a signed and normalized loss, deterministic LLM decoding, a sound unit system, a deterministic MI pipeline, a duplicate-aware population schedule with a provably correct 2D frontier, an MCTS explainer, a synthetic protocol, and strict acceptance criteria. The design is implementation-ready and empirically grounded: the behavior I report in results follows from these contracts rather than ad hoc choices.